

1. Problem:

In [174]:

```
1 import numpy as np
2 import pandas as pd
3 import sqlalchemy as sa
4
5 from statsmodels.stats.outliers_influence import variac
6 from sklearn.linear_model import LogisticRegression
7 from sklearn.model_selection import train_test_split,
8
9 from sklearn.metrics import accuracy_score, f1_score, re
10 from sklearn.metrics import classification_report, co
11 from sklearn.metrics import roc_curve
12 from sklearn.metrics import auc
13
14 from imblearn.over_sampling import SMOTE
15 from imblearn.over_sampling import RandomOverSampler
16
17 from summarytools import dfSummary
18 import seaborn as sns
19 import matplotlib.pyplot as plt
20 import plotly.express as px
21 import pickle
22
23 import warnings
24 warnings.filterwarnings('ignore')
```

2. Data Gathering

In []:

```
1  
2  
3 This text was recognized by the built-in Ocrad engine.
```

In [22]:

```
1 # connecting mySQL database to Jupyter Notebook for da  
2 con = sa.create_engine("mysql+pymysql://root:@Localhos  
3 con
```

Out[22]:

```
Engine(mysql+pymysql://root:***@Localhost:3  
306/diabetes_db)
```

In [23]:

```
1 df = pd.read_sql_table('diabetes', con)
2 df
```

Out[23]:

	Glucose	BloodPressure	SkinThickness	Insulin	
0	148	50	35	0	...
1	85	66	29	0	...
2	183	64	0	0	...
3	150	66	23	94	...
4	150	40	35	168	...
...
763	101	76	48	180	...
764	122	70	27	0	...
765	121	72	23	112	...
766	126	60	0	0	...
767	93	70	31	0	...

768 rows × 8 columns



In [24]:

```
1 df.head()
```

Out[24]:

	Glucose	BloodPressure	SkinThickness	Insulin	BM
0	148	50	35	0	33.
1	85	66	29	0	26.
2	183	64	0	0	23.
3	150	66	23	94	28.
4	150	40	35	168	43.

3. Exploratory Data Analysis

In [25]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 8 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Glucose                768 non-null   int64
 1   BloodPressure          768 non-null   int64
 2   SkinThickness          768 non-null   int64
 3   Insulin                768 non-null   int64
 4   BMI                    768 non-null   float64
 5   DiabetesPedigreeFunction 768 non-null   float64
 6   Age                    768 non-null   int64
 7   Outcome                768 non-null   int64
dtypes: float64(2), int64(6)
memory usage: 48.1 KB
```



In [26]:




```
1 dfSummary(df)
```




Out[26]:

Data Frame Summary

Dimensions: 768 x 8
Duplicates: 0

No	Variable	Stats / Values	Freqs / (% of Valid)	Graph
1	Glucose [int64]	Mean		
		(sd) :		
		121.1		
		(31.8)		
		min <		
		med <		
		max:	136	
		0.0 <	distinct	
		117.0	values	
		<		
		199.0		
		IQR		
2	BloodPressure [int64]	(CV) :		
		43.0		
		(3.8)		
		Mean		
		(sd) :		
		69.1		
		(19.4)		
		min <		
		med <		
		max:	47	
		0.0 <	distinct	
		72.0 <	values	
		122.0		
		IQR		
		(CV) :		
		18.0		
		(3.6)		

No	Variable	Stats / Values	Freqs / (% of Valid)	Graph
3	SkinThickness [int64]	Mean		
		(sd) :		
		20.5		
		(16.0)		
		min <		
		med <		
		max:	51	
		0.0 <	distinct	
		23.0 <	values	
		99.0		
		IQR		
		(CV) :		
		32.0		
		(1.3)		
4	Insulin [int64]	Mean		
		(sd) :		
		79.8		
		(115.2)		
		min <		
		med <		
		max:	186	
		0.0 <	distinct	
		30.5 <	values	
		846.0		
		IQR		
		(CV) :		
		127.2		
		(0.7)		
5	BMI [float64]	Mean		
		(sd) :		
		32.0		
		(7.9)		
		min <		
		med <		
		max:	248	
		0.0 <	distinct	
		32.0 <	values	
		67.1		
		IQR		
		(CV) :		
		9.3		
		(4.1)		

No	Variable	Stats / Values	Freqs / (% of Valid)	Graph
6	DiabetesPedigreeFunction [float64]	Mean		
		(sd) :		
		0.5		
		(0.3)		
		min <		
		med <		
		max:	517	
		0.1 <	distinct	
		0.4 <	values	
		2.4		
		IQR		
		(CV) :		
		0.4		
		(1.4)		
7	Age [int64]	Mean		
		(sd) :		
		33.2		
		(11.8)		
		min <		
		med <		
		max:	52	
		21.0 <	distinct	
		29.0 <	values	
		81.0		
		IQR		
		(CV) :		
		17.0		
		(2.8)		
In [27]:		Mean		
		(sd) :		
1	# x = df.drop('Outcome', axis=1)	0.5		
2	# y = df['Outcome']	(0.5)		
3		min <		
4	# x_train,x_test,y_train,y_test= train_test_split(x,	med <		
8	Outcome	max:	2	
	[int64]	0.0 <	distinct	
		0.0 <	values	
		1.0		
In [28]:		IQR		
1	# x_train	(CV) :		
		1.0		
		(0.7)		

1. Glucose

In [29]:

```
1 df['Glucose'].head()
```

Out[29]:

```
0    148
1     85
2    183
3    150
4    150
```

Name: Glucose, dtype: int64

In [30]:

```
1 df['Glucose'].info()
```

```
<class 'pandas.core.series.Series'>
RangeIndex: 768 entries, 0 to 767
Series name: Glucose
Non-Null Count  Dtype
-----
768 non-null    int64
dtypes: int64(1)
memory usage: 6.1 KB
```

In [31]:

```
1 df['Glucose'].isna().sum()
```

Out[31]:

```
0
```

In [32]:

```
1 df['Glucose'].value_counts()
```

Out[32]:

```
100    17
99     17
150    15
106    14
129    14
..
44     1
177    1
191    1
61     1
190    1
```


Name: Glucose, Length: 136, dtype: int64

In [40]:

```
1 plt.figure(figsize= (30,6))
2 sns.countplot(x=df["Glucose"])
```

Out[40]:

<AxesSubplot:xlabel='Glucose', ylabel='count'>




In [41]:

```
1 plt.figure(figsize=(15,5))  
2 sns.boxplot(df['Glucose'])
```

Out[41]:

<AxesSubplot: xlabel='Glucose'>



2. BloodPressure

In [44]:

```
1 df['BloodPressure'].head()
```

Out[44]:

```
0    50  
1    66  
2    64  
3    66  
4    40
```

Name: BloodPressure, dtype: int64

In [45]:

```
1 df['BloodPressure'].info()
```

```
<class 'pandas.core.series.Series'>  
RangeIndex: 768 entries, 0 to 767  
Series name: BloodPressure  
Non-Null Count  Dtype  
-----  
768 non-null    int64  
dtypes: int64(1)  
memory usage: 6.1 KB
```

In [47]:

```
1 df['BloodPressure'].isna().sum()
```

Out[47]:

0

In [48]:

```
1 df['BloodPressure'].value_counts()
```

Out[48]:

```
70      57
74      52
78      45
68      45
64      43
72      43
80      40
76      39
60      37
```


```
In [50]:35
```

```
plt.figure(figsize= (30,6))
sns.countplot(df["BloodPressure"])
```

```
Out[50]:25
```

```
23
```

```
plt.xlabel='BloodPressure', ylabel
count
```



```
Out[50]:25
```

```
23
```


```
In [51]:7
```

```
plt.figure(figsize= (15,5))
sns.boxplot(df["BloodPressure"])
```

```
Out[51]:4
```

```
4
```

```
plt.xlabel='BloodPressure'>
```



```
Out[51]:4
```

```
4
```



```
102      1
```

```
61      1
```

```
24      1
```

BloodPressure

```
38      1
```

```
40 [53]:1
```

```
114     1
```

```
Name: df['SkinThickness'].head()
Name: BloodPressure, dtype: int64
```

Out[53]:

```
0      35
```

```
1      29
```

```
2       0
```

```
3      23
```

```
4      35
```

```
Name: SkinThickness, dtype: int64
```

In [54]:

```
1 df['SkinThickness'].info()
```

```
<class 'pandas.core.series.Series'>
```

```
RangeIndex: 768 entries, 0 to 767
```

```
Series name: SkinThickness
```

```
Non-Null Count  Dtype
```

```
-----
```

```
768 non-null    int64
```

```
dtypes: int64(1)
```

```
memory usage: 6.1 KB
```

In [56]:

```
1 df['SkinThickness'].isna().sum()
```

Out[56]:

```
0
```

In [55]:

```
1 df['SkinThickness'].value_counts()
```

Out[55]:


0 227
32 31
30 27
27 23
23 22
33 20
28 20
18 20
31 19
19 18
39 18
29 17
40 16

25 [57] 16

```
26 16  
22 16 plt.figure(figsize= (30,6))  
37 16 sns.countplot(df["SkinThickness"])
```

41 15
Out[57]: 15

36 14
<AxesSubplot:xlabel='SkinThickness', ylabel
15 count>
17 14




41 15
46 8
34 8
12 7
38 7
11 6
43 6
16 6
45 6
14 6
44 5
10 5
48 4
47 4
49 3

```
50      3
In [58]:
```

```
7      2
52     1 plt.figure(figsize= (15,5))
54     2 sns.boxplot(df["SkinThickness"])
```

```
63     1
Out[58]:
```

```
60     1
56     1 AxesSubplot: xlabel='SkinThickness'>
51     1
```



4. Insulin

```
In [59]:
```

```
1 df['Insulin'].head()
```

```
Out[59]:
```

```
0      0
1      0
2      0
3     94
4    168
```

```
Name: Insulin, dtype: int64
```

In [60]:

```
1 df['Insulin'].info()
```

```
<class 'pandas.core.series.Series'>  
RangeIndex: 768 entries, 0 to 767  
Series name: Insulin  
Non-Null Count  Dtype  
-----  
768 non-null    int64  
dtypes: int64(1)  
memory usage: 6.1 KB
```

In [61]:

```
1 df['Insulin'].isna().sum()
```

Out[61]:

0

In [62]:

```
1 df['Insulin'].value_counts()
```

Out[62]:

```
0      374  
105     11  
130      9  
140      9  
120      8  
...  
73       1  
171      1  
255      1  
52       1  
112      1
```


Name: Insulin, Length: 186, dtype: int64

In [63]:

```
1 plt.figure(figsize= (30,6))
2 sns.countplot(df["Insulin"])
```

Out[63]:

<AxesSubplot:xlabel='Insulin', ylabel='count'>




In [64]:

```
1 plt.figure(figsize= (15,5))
2 sns.boxplot(df["Insulin"])
```

Out[64]:

<AxesSubplot:xlabel='Insulin'>



5. BMI

In [65]:

```
1 df['BMI'].head()
```

Out[65]:

```
0    33.6
1    26.6
2    23.3
3    28.1
4    43.1
```

Name: BMI, dtype: float64

In [66]:

```
1 df['BMI'].info()
```

```
<class 'pandas.core.series.Series'>
RangeIndex: 768 entries, 0 to 767
Series name: BMI
Non-Null Count  Dtype
-----
768 non-null    float64
dtypes: float64(1)
memory usage: 6.1 KB
```

In [67]:

```
1 df['BMI'].isna().sum()
```

Out[67]:

```
0
```


In [68]:

```
1 df['BMI'].value_counts()
```

Out[68]:

```
32.0    13
31.6    12
31.2    12
0.0     11
32.4    10
..
36.7     1
41.8     1
42.6     1
42.8     1
46.3     1
```


Name: BMI, Length: 248, dtype: int64

In [70]:

```
1 plt.figure(figsize= (30,6))
2 sns.countplot(df["BMI"])
```

Out[70]:

<AxesSubplot:xlabel='BMI', ylabel='count'>




In [71]:

```
1 plt.figure(figsize= (15,5))
2 sns.boxplot(df["BMI"])
```

Out[71]:

<AxesSubplot:xlabel='BMI'>



6. DiabetesPedigreeFunction

In [72]:

```
1 df['DiabetesPedigreeFunction'].head()
```

Out[72]:

```
0    0.627
1    0.351
2    0.672
3    0.167
4    2.288
```

Name: DiabetesPedigreeFunction, dtype: float64

In [73]:

```
1 df['DiabetesPedigreeFunction'].info()
```

```
<class 'pandas.core.series.Series'>
RangeIndex: 768 entries, 0 to 767
Series name: DiabetesPedigreeFunction
Non-Null Count  Dtype
-----
768 non-null    float64
dtypes: float64(1)
memory usage: 6.1 KB
```

In [74]:

```
1 df['DiabetesPedigreeFunction'].isna().sum()
```

Out[74]:

0

In [75]:

```
1 df['DiabetesPedigreeFunction'].value_counts()
```

Out[75]:

```
0.258    6
0.254    6
0.268    5
0.207    5
0.261    5
..
1.353    1
0.655    1
0.092    1
0.926    1
0.171    1
```


```
Name: DiabetesPedigreeFunction, Length: 51
7, dtype: int64
```

In [78]:

```
1 plt.figure(figsize= (30,6))
2 sns.countplot(df[ "DiabetesPedigreeFunction"])
```

Out[78]:

<AxesSubplot:xlabel='DiabetesPedigreeFunction', ylabel='count'>




In [77]:

```
1 plt.figure(figsize= (15,5))
2 sns.boxplot(df[ "DiabetesPedigreeFunction"])
```

Out[77]:

<AxesSubplot:xlabel='DiabetesPedigreeFunction', ylabel='count'>



7. Age

In [80]:

```
1 df['Age'].head()
```

Out[80]:

```
0    50
1    31
2    52
3    21
4    33
```

Name: Age, dtype: int64

In [81]:

```
1 df['Age'].info()
```

```
<class 'pandas.core.series.Series'>
RangeIndex: 768 entries, 0 to 767
Series name: Age
Non-Null Count  Dtype
-----
768 non-null    int64
dtypes: int64(1)
memory usage: 6.1 KB
```

In [82]:

```
1 df['Age'].isna().sum()
```

Out[82]:

```
0
```

In [83]:

```
1 df['Age'].value_counts()
```

Out[83]:

```
22 72
21 63
25 48
24 46
23 38
28 35
26 33
27 32
29 29
31 24
41 22
30 21
37 19
42 18
```


```
33 [85]:
```

```
38 16
36 16 plt.figure(figsize= (30,6))
32 15 sns.countplot(df["Age"])
```

```
45 15
```

```
Out[85]:
```

```
34 14
43 13
<AxesSubplot:xlabel='Age', ylabel='count'>
46 13
```



```
50 8
51 8
44 8
58 7
47 6
54 6
49 5
48 5
57 5
60 5
66 4
53 4
62 4
55 4
63 4
```



```


67     3
In [86]:
58     3
59     1 plt.figure(figsize= (15,5))
65     3 sns.boxplot(df[ "Age" ])
69     2

```

```

61     2
Out[86]:
72     1
81     1 AxesSubplot: xlabel='Age' >
64     1

```



8. Outcome

In [89]:

```

1 df[ 'Outcome' ]

```

Out[89]:

```

0     1
1     0
2     1
3     0
4     1
..
763    0
764    0
765    0
766    1
767    0

```

Name: Outcome, Length: 768, dtype: int64

In [90]:

```
1 df['Outcome'].info()
```

```
<class 'pandas.core.series.Series'>  
RangeIndex: 768 entries, 0 to 767  
Series name: Outcome  
Non-Null Count  Dtype  
-----  
768 non-null    int64  
dtypes: int64(1)  
memory usage: 6.1 KB
```

In [91]:

```
1 df['Outcome'].isna().sum()
```

Out[91]:


0

In [94]:

```
1 sns.countplot(df['Outcome'])
```

Out[94]:

<AxesSubplot:xlabel='Outcome', ylabel='count'>




In [95]:

```
1 plt.figure(figsize= (15,5))  
2 sns.boxplot(df[ "Outcome"])
```

Out[95]:

<AxesSubplot:xlabel='Outcome'>



4. Feature Engineering

In [97]:

```
1 sns.pairplot(df, hue = 'Outcome')
```

Out[97]:

<seaborn.axisgrid.PairGrid at 0x198b6bdba30

>



4.1 zero value Imputation

In [113]:

```
1 num= df[df["SkinThickness"]==0]
2 num1= df[df["BloodPressure"]==0]
3 num2= df[df["Glucose"]==0]
4 num3= df[df["Insulin"]==0]
5 num4= df[df["BMI"]==0]
6 num.shape,num1.shape,num2.shape,num3.shape,num4.shape
```

Out[113]:

```
((227, 8), (35, 8), (5, 8), (374, 8), (11, 8))
```

In [114]:

```
1 df[['Glucose', 'BloodPressure', 'BMI', 'Insulin', 'SkinThickness']]
```

In [115]:

```
1 df.isna().sum()
```

Out[115]:

Glucose	5
BloodPressure	35
SkinThickness	227
Insulin	374
BMI	11
DiabetesPedigreeFunction	0
Age	0
Outcome	0
dtype:	int64

In [117]:

```
1 df[['Glucose', 'BloodPressure', 'BMI', 'Insulin', 'SkinT
```

In [119]:

```
1 num= df[df["SkinThickness"]==0]
2 num1= df[df["BloodPressure"]==0]
3 num2= df[df["Glucose"]==0]
4 num3= df[df["Insulin"]==0]
5 num4= df[df["BMI"]==0]
6 num.shape, num1.shape, num2.shape, num3.shape, num4.shape
```

Out[119]:

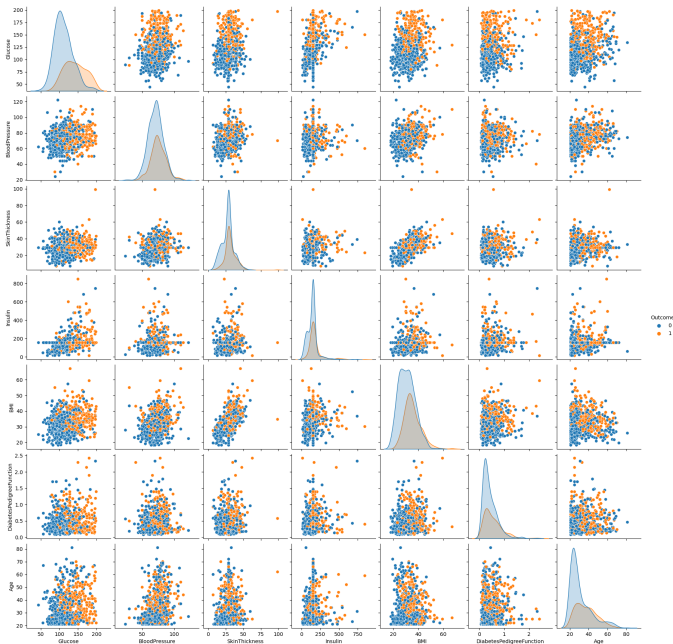
```
((0, 8), (0, 8), (0, 8), (0, 8), (0, 8))
```

In [132]:

```
1 sns.pairplot(df, hue = 'Outcome')
```

Out[132]:

<seaborn.axisgrid.PairGrid at 0x198bbefec40>



4.2 Outlier Imputation

In [134]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 8 columns):
 #   Column                Non-Null Cou
nt  Dtype                nt
---  ---
0   Glucose                768 non-null
float64
1   BloodPressure          768 non-null
float64
2   SkinThickness          768 non-null
float64
3   Insulin                768 non-null
float64
4   BMI                    768 non-null
float64
5   DiabetesPedigreeFunction 768 non-null
float64
6   Age                    768 non-null
int64
7   Outcome                768 non-null
int64
dtypes: float64(6), int64(2)
memory usage: 48.1 KB
```


1.Glucose

In [136]:

```
1 plt.figure(figsize= (15,5))
2 sns.boxplot(df["Glucose"])
```

Out[136]:

<AxesSubplot:xlabel='Glucose'>




2. BloodPressure

In [137]:

```
1 plt.figure(figsize= (15,5))
2 sns.boxplot(df["BloodPressure"])
```

Out[137]:

<AxesSubplot:xlabel='BloodPressure'>




In [181]:

```
1  # Removing outliers with IQR method
2
3  q1 = df['BloodPressure'].quantile(0.25)
4  q2 = df['BloodPressure'].quantile(0.50)
5  q3 = df['BloodPressure'].quantile(0.75)
6  median = df['BloodPressure'].median()
7
8  iqr = q3 - q1
9
10 upper_tail = q3 + 1.5 * iqr
11 lower_tail = q1 - 1.5 * iqr
12
13 print("Q1 :", q1)
14 print("Q2 :", q2)
15 print("Q3 :", q3)
16 print("Median :", median)
17
18 print("upper_tail :", upper_tail)
19 print("lower_tail :", lower_tail)
20 a = df['BloodPressure'].loc[(df['BloodPressure'] > upper_tail)]
21 print("upper_tail outliers:\n ", a)
22 b = df['BloodPressure'].loc[(df['BloodPressure'] < lower_tail)]
23 print("lower_tail outliers: \n", b)
24 df.loc[(df['BloodPressure'] > upper_tail), 'BloodPressure'] = upper_tail
25 df.loc[(df['BloodPressure'] < lower_tail), 'BloodPressure'] = lower_tail
26
27 plt.figure(figsize=(15,3))
28 sns.boxplot(df['BloodPressure'])
```

```
Q1 : 64.0
Q2 : 72.18758526603001
Q3 : 80.0
Median : 72.18758526603001
upper_tail : 104.0
lower_tail : 40.0
upper_tail outliers:
Series([], Name: BloodPressure, dtype: float64)
lower_tail outliers:
18      30.0
125     30.0
597     24.0
599     38.0
Name: BloodPressure, dtype: float64
```

Out[181]:

<AxesSubplot:xlabel='BloodPressure'>




3. SkinThickness

In [138]:

```
1 plt.figure(figsize= (15,5))  
2 sns.boxplot(df["SkinThickness"])
```

Out[138]:

<AxesSubplot:xlabel='SkinThickness'>



In [184]:

```
1  # Removing outliers with IQR method
2
3  q1 = df['SkinThickness'].quantile(0.25)
4  q2 = df['SkinThickness'].quantile(0.50)
5  q3 = df['SkinThickness'].quantile(0.75)
6  median = df['SkinThickness'].median()
7
8  iqr = q3 - q1
9
10 upper_tail = q3 + 1.5 * iqr
11 lower_tail = q1 - 1.5 * iqr
12
13 print("Q1 :", q1)
14 print("Q2 :", q2)
15 print("Q3 :", q3)
16 print("Median :", median)
17
18 print("upper_tail :", upper_tail)
19 print("lower_tail :", lower_tail)
20 a = df['SkinThickness'].loc[(df['SkinThickness'] > upper_tail)]
21 print("upper_tail outliers:\n ", a)
22 b = df['SkinThickness'].loc[(df['SkinThickness'] < lower_tail)]
23 print("lower_tail outliers: \n", b)
24 df.loc[(df['SkinThickness'] > upper_tail), 'SkinThickness'] = df['SkinThickness'].loc[(df['SkinThickness'] > upper_tail)]
25 df.loc[(df['SkinThickness'] < lower_tail), 'SkinThickness'] = df['SkinThickness'].loc[(df['SkinThickness'] < lower_tail)]
26
```


```
Q1 : 25.0
Q2 : 29.153419593345657
Q3 : 32.0
Median : 29.153419593345657
upper_tail : 42.5
lower_tail : 14.5
upper_tail outliers:
Series([], Name: SkinThickness, dtype: float64)
lower_tail outliers:
Series([], Name: SkinThickness, dtype: float64)
```

In [185]:

```
1 plt.figure(figsize=(15,3))
2 sns.boxplot(df['SkinThickness'])
```

Out[185]:

<AxesSubplot: xlabel='SkinThickness'>




4. Insulin

In [139]:

```
1 plt.figure(figsize= (15,5))  
2 sns.boxplot(df["Insulin"])
```

Out[139]:

<AxesSubplot:xlabel='Insulin'>



In [187]:

```
1  # Removing outliers with IQR method
2
3  q1 = df['Insulin'].quantile(0.25)
4  q2 = df['Insulin'].quantile(0.50)
5  q3 = df['Insulin'].quantile(0.75)
6  median = df['Insulin'].median()
7
8  iqr = q3 - q1
9
10 upper_tail = q3 + 1.5 * iqr
11 lower_tail = q1 - 1.5 * iqr
12
13 print("Q1 :", q1)
14 print("Q2 :", q2)
15 print("Q3 :", q3)
16 print("Median :", median)
17
18 print("upper_tail :", upper_tail)
19 print("lower_tail :", lower_tail)
20 a = df['Insulin'].loc[(df['Insulin'] > upper_tail)]
21 print("upper_tail outliers:\n ", a)
22 b = df['Insulin'].loc[(df['Insulin'] < lower_tail)]
23 print("lower_tail outliers: \n", b)
24 df.loc[(df['Insulin'] > upper_tail), 'Insulin'] = upper_tail
25 df.loc[(df['Insulin'] < lower_tail), 'Insulin'] = lower_tail
26
```


```
Q1 : 121.5
Q2 : 155.5482233502538
Q3 : 155.5482233502538
Median : 155.5482233502538
upper_tail : 206.62055837563452
lower_tail : 70.42766497461929
upper_tail outliers:
      8      543.0
13      846.0
16      230.0
20      235.0
31      245.0
...
707     335.0
710     387.0
713     291.0
715     392.0
753     510.0
Name: Insulin, Length: 82, dtype: float64
lower_tail outliers:
      32      54.0
40      70.0
51      36.0
52      23.0
68      38.0
...
672     49.0
680     45.0
711     22.0
747     57.0
760     16.0
Name: Insulin, Length: 82, dtype: float64
```

In [188]:

```
1 plt.figure(figsize= (15,5))
2 sns.boxplot(df["Insulin"])
```

Out[188]:

<AxesSubplot:xlabel='Insulin'>




5. BMI

In [141]:

```
1 plt.figure(figsize= (15,5))
2 sns.boxplot(df["BMI"])
```

Out[141]:

<AxesSubplot:xlabel='BMI'>



In [189]:

```
1  # Removing outliers with IQR method
2
3  q1 = df['BMI'].quantile(0.25)
4  q2 = df['BMI'].quantile(0.50)
5  q3 = df['BMI'].quantile(0.75)
6  median = df['BMI'].median()
7
8  iqr = q3 - q1
9
10 upper_tail = q3 + 1.5 * iqr
11 lower_tail = q1 - 1.5 * iqr
12
13 print("Q1 :", q1)
14 print("Q2 :", q2)
15 print("Q3 :", q3)
16 print("Median :", median)
17
18 print("upper_tail :", upper_tail)
19 print("lower_tail :", lower_tail)
20 a = df['BMI'].loc[(df['BMI'] > upper_tail)]
21 print("upper_tail outliers:\n ", a)
22 b = df['BMI'].loc[(df['BMI'] < lower_tail)]
23 print("lower_tail outliers: \n", b)
24 df.loc[(df['BMI'] > upper_tail), 'BMI'] = upper_tail
25 df.loc[(df['BMI'] < lower_tail), 'BMI'] = lower_tail
26
```


```
Q1 : 27.5
Q2 : 32.4
Q3 : 36.6
Median : 32.4
upper_tail : 50.25
lower_tail : 13.849999999999998
upper_tail outliers:
    120    53.2
    125    55.0
    177    67.1
    193    52.3
    247    52.3
    303    52.9
    445    59.4
    673    57.3
Name: BMI, dtype: float64
lower_tail outliers:
    Series([], Name: BMI, dtype: float64)
```

In [190]:

```
1 plt.figure(figsize= (15,5))
2 sns.boxplot(df["BMI"])
```

Out[190]:

<AxesSubplot:xlabel='BMI'>




6. DiabetesPedigreeFunction

In [142]:

```
1 plt.figure(figsize= (15,5))
2 sns.boxplot(df[ "DiabetesPedigreeFunction"])
```

Out[142]:

<AxesSubplot:xlabel='DiabetesPedigreeFunction'>



In [191]:

```
1  # Removing outliers with IQR method
2
3  q1 = df['DiabetesPedigreeFunction'].quantile(0.25)
4  q2 = df['DiabetesPedigreeFunction'].quantile(0.50)
5  q3 = df['DiabetesPedigreeFunction'].quantile(0.75)
6  median = df['DiabetesPedigreeFunction'].median()
7
8  iqr = q3 - q1
9
10 upper_tail = q3 + 1.5 * iqr
11 lower_tail = q1 - 1.5 * iqr
12
13 print("Q1 :", q1)
14 print("Q2 :", q2)
15 print("Q3 :", q3)
16 print("Median :", median)
17
18 print("upper_tail :", upper_tail)
19 print("lower_tail :", lower_tail)
20 a = df['DiabetesPedigreeFunction'].loc[(df['DiabetesPe
21 print("upper_tail outliers:\n ", a)
22 b= df['DiabetesPedigreeFunction'].loc[(df['DiabetesPec
23 print("lower_tail outliers: \n", b)
24 df.loc[(df['DiabetesPedigreeFunction']> upper_tail), 'I
25 df.loc[(df['DiabetesPedigreeFunction']< lower_tail), 'I
26
```


```
Q1 : 0.24375
Q2 : 0.3725
Q3 : 0.62625
Median : 0.3725
upper_tail : 1.2
lower_tail : -0.32999999999999996
upper_tail outliers:
  4      2.288
12      1.441
39      1.390
45      1.893
58      1.781
100     1.222
147     1.400
187     1.321
218     1.224
228     2.329
243     1.318
245     1.213
259     1.353
292     1.224
308     1.391
330     1.476
370     2.137
371     1.731
383     1.268
395     1.600
445     2.420
534     1.251
593     1.699
606     1.258
618     1.282
621     1.698
622     1.461
659     1.292
661     1.394
Name: DiabetesPedigreeFunction, dtype: float64
lower_tail outliers:
  Series([], Name: DiabetesPedigreeFunction,
  dtype: float64)
```


In [192]:

```
1 plt.figure(figsize= (15,5))
2 sns.boxplot(df["DiabetesPedigreeFunction"])
```

Out[192]:

<AxesSubplot:xlabel='DiabetesPedigreeFunction'>




7. Age

In [143]:

```
1 plt.figure(figsize= (15,5))
2 sns.boxplot(df["Age"])
```

Out[143]:

<AxesSubplot:xlabel='Age'>



In [193]:

```
1  # Removing outliers with IQR method
2
3  q1 = df['Age'].quantile(0.25)
4  q2 = df['Age'].quantile(0.50)
5  q3 = df['Age'].quantile(0.75)
6  median = df['Age'].median()
7
8  iqr = q3 - q1
9
10 upper_tail = q3 + 1.5 * iqr
11 lower_tail = q1 - 1.5 * iqr
12
13 print("Q1 :", q1)
14 print("Q2 :", q2)
15 print("Q3 :", q3)
16 print("Median :", median)
17
18 print("upper_tail :", upper_tail)
19 print("lower_tail :", lower_tail)
20 a = df['Age'].loc[(df['Age'] > upper_tail)]
21 print("upper_tail outliers:\n ", a)
22 b = df['Age'].loc[(df['Age'] < lower_tail)]
23 print("lower_tail outliers: \n", b)
24 df.loc[(df['Age'] > upper_tail), 'Age'] = upper_tail
25 df.loc[(df['Age'] < lower_tail), 'Age'] = lower_tail
26
```


```
Q1 : 24.0
Q2 : 29.0
Q3 : 41.0
Median : 29.0
upper_tail : 66.5
lower_tail : -1.5
upper_tail outliers:
    123    69
363    67
453    72
459    81
489    67
537    67
666    70
674    68
684    69
Name: Age, dtype: int64
lower_tail outliers:
Series([], Name: Age, dtype: int64)
```

In [194]:

```
1 plt.figure(figsize= (15,5))
2 sns.boxplot(df["Age"])
```

Out[194]:

<AxesSubplot: xlabel='Age'>



5. Feature Selection

5.1 Assumption

In [144]:

```
1 df.corr()
```

Out[144]:


	Glucose	BloodPressure	Sk
Glucose	1.000000	0.215938	
BloodPressure	0.215938	1.000000	
SkinThickness	0.192013	0.190815	
Insulin	0.395345	0.072356	
BMI	0.232153	0.280253	
DiabetesPedigreeFunction	0.134736	-0.003836	
Age	0.259389	0.319440	
Outcome	0.489601	0.162449	

In [145]:

```
1 plt.figure(figsize=(15,5))
2 sns.heatmap(df.corr(), annot=True)
```

Out[145]:

<AxesSubplot:>



5.2. No Multicollinearity

In [148]:

```
1 df1 = df.drop('Outcome', axis= 1)
2 df1.head()
```

Out[148]:

	Glucose	BloodPressure	SkinThickness	Insulin
0	148.0	50.0	35.00000	155.548223
1	85.0	66.0	29.00000	155.548223
2	183.0	64.0	29.15342	155.548223
3	150.0	66.0	23.00000	94.000000
4	150.0	40.0	35.00000	168.000000

In [151]:

```
1 vif_data = pd.DataFrame()
2 vif_data["feature"] = df1.columns
3
4 # calculating VIF for each feature
5 vif_data["VIF"] = [variance_inflation_factor(df1.values,
6     for i in range(len(df1.columns)))
7
8 print(vif_data)
```

	feature	VIF
0	Glucose	21.347959
1	BloodPressure	31.335650
2	SkinThickness	17.360263
3	Insulin	5.232267
4	BMI	33.628317
5	DiabetesPedigreeFunction	3.131796
6	Age	10.680453

6. Model Building

6.1. sampling

In []:

```
1
```

6.2. Splitting the data into Training data and Testing data

In [227]:

```
1 # regular Model
2 # independent variable(x)
3 x = df.drop('Outcome', axis= 1)
4 # dependent variable(y)
5 y = df['Outcome']
6 x_train, x_test, y_train, y_test = train_test_split(x,
```

In [228]:

```
1 print("shape of x_train:",x_train.shape)
2 print("shape of y_train:",y_train.shape)
3 print("shape of x_test:",x_test.shape)
4 print("shape of y_test:",y_test.shape)
```

shape of x_train: (537, 7)

shape of y_train: (537,)

shape of x_test: (231, 7)

shape of y_test: (231,)

6.3 Model Training

Logistic Regression

In [229]:

```
1 logistic_model = LogisticRegression()  
2 logistic_model.fit(x_train,y_train)
```

Out[229]:

LogisticRegression()

In [230]:

```
1 logistic_pred =logistic_model.predict(x_test)  
2 logistic_pred[:5]
```

Out[230]:

array([0, 0, 1, 0, 0], dtype=int64)

6.4 Model Evaluation

Fucntion for test and train

In [231]:

```
1  # Testing Data
2  def testing_evaluation(model,x,y):
3      prediction = model.predict(x)
4
5      cnf_matrix = confusion_matrix(y, prediction)
6      print("Confusion Matrix:\n", cnf_matrix)
7      print(""*45)
8
9      accuracy = accuracy_score(y, prediction)
10     print('Accuracy \n',accuracy)
11     print(""*45)
12
13     clf_report= classification_report(y, prediction)
14     print("Classification Report\n", clf_report)
```

In [232]:

```
1  # Training Data
2  def training_evaluation(model,x,y):
3      prediction = model.predict(x)
4
5      cnf_matrix = confusion_matrix(y, prediction)
6      print("Confusion Matrix:\n", cnf_matrix)
7      print(""*45)
8
9      accuracy = accuracy_score(y, prediction)
10     print('Accuracy \n',accuracy)
11     print(""*45)
12
13     clf_report= classification_report(y, prediction)
14     print("Classification Report\n", clf_report)
```

regular Model

In [233]:

```
1 # Testing Data
2 testing_evaluation(logistic_model,x_test,y_test)
```

Confusion Matrix:

```
[[135  15]
 [ 35  46]]
*****
**
```

Accuracy

```
0.7835497835497836
*****
**
```

Classification Report


		precision	recall	f1-scor
e	support			
	0	0.79	0.90	0.84
150				
	1	0.75	0.57	0.65
81				
	accuracy			0.78
231				
	macro avg	0.77	0.73	0.75
231				
	weighted avg	0.78	0.78	0.78
231				

In [234]:

```
1 plot_confusion_matrix(logistic_model,x_test,y_test)
```

Out[234]:

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x198c466ae50>
```



In [235]:

```
1 # Training Data
2 training_evaluation(logistic_model,x_train,y_train)
```

Confusion Matrix:

```
[[309  41]
 [ 84 103]]
```

**

Accuracy

0.7672253258845437

**

Classification Report


		precision	recall	f1-scor
e	support			
	0	0.79	0.88	0.83
350				
	1	0.72	0.55	0.62
187				
	accuracy			0.77
537				
	macro avg	0.75	0.72	0.73
537				
	weighted avg	0.76	0.77	0.76
537				

In [236]:

```
1 plot_confusion_matrix(logistic_model,x_train,y_train)
```

Out[236]:

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x198c7033d60>
```



checking roc curve :

In [237]:


```
1 y_pred_proba = logistic_model.predict_proba(x_train)
2 y_pred_proba[:,1]
3 fpr,tpr,thresh = roc_curve(y_train,y_pred_proba[:,1])
```

In [238]:

```
1 plt.plot(fpr, tpr)
2 plt.xlabel("False Positive Rate Value(FPR)")
3 plt.ylabel("True Positive Rate (TPR)")
4 plt.title("ROC curve")
```

Out[238]:

Text(0.5, 1.0, 'ROC curve')



In [239]:

```
1 ### auc value
2 auc_value = auc(fpr, tpr)
3 auc_value
```

Out[239]:

0.8310313216195568

In [240]:

```
1 thresh = np.arange(0,1,0.1)
2 info_df = pd.DataFrame()
3
4 for i in thresh:
5     preds = (logistic_model.predict_proba(x_test)[:,-1])
6
7     thresh_df=pd.DataFrame(data=[accuracy_score(y_test, preds),
8                                     recall_score(y_test, preds)],
9                             index=["Accuracy", "Recall"],
10                                columns=[i])
11     info_df = pd.concat([info_df,thresh_df],axis=1)
12
13 info_df.columns=thresh
14 info_df
```

Out[240]:

	0.0	0.1	0.2	0.3	
Accuracy	0.350649	0.571429	0.718615	0.774892	0.78
Recall	1.000000	1.000000	0.913580	0.765432	0.66
Precision	0.350649	0.450000	0.560606	0.652632	0.71
F1-score	0.519231	0.620690	0.694836	0.704545	0.68

Hyper-Parameter Tunning :

In [241]:

```
1 weights = np.linspace(0,0.99,num=100)
2 param_grid = {"C":np.arange(0.1,20),"penalty":["l1","l2"]}
3
4
5 gscv_model = GridSearchCV(logistic_model,param_grid)
6 gscv_model.fit(x_train,y_train)
```

Out[241]:


```

GridSearchCV(estimator=LogisticRegression
()),
               param_grid={'C': array([ 0.1,
1.1, 2.1, 3.1, 4.1, 5.1, 6.1, 7.1,
8.1, 9.1, 10.1,
               11.1, 12.1, 13.1, 14.1, 15.1, 16.1,
17.1, 18.1, 19.1]),
               'class_weight':
[{'0': 0.0, 1: 1.0}, {'0': 0.01, 1: 0.99},
{'0': 0.02, 1: 0.98},
{'0': 0.03, 1: 0.97},
{'0': 0.04, 1: 0.96},
{'0': 0.05, 1: 0.95},
{'0': 0.06, 1: 0.94},
{'0': 0.07, 1: 0.9299999999999999},
{'0': 0.08, 1: 0.92},
{'0': 0.09, 1: 0.91}, {'0':...: 0.9},
{'0': 0.11, 1: 0.89}],
               'penalty': 'l2'})

In [409]: gscv_model.best_params_
Out[409]: {'C': 0.11, 'class_weight': {0: 0.89, 1: 0.85}, 'penalty': 'l2'}

In [342]: gscv_model.best_params_
Out[342]: {'C': 0.12, 'class_weight': {0: 0.88, 1: 0.85}, 'penalty': 'l2'}

In [343]: gscv_model.best_params_
Out[343]: {'C': 0.13, 'class_weight': {0: 0.87, 1: 0.84}, 'penalty': 'l2'}

In [344]: gscv_model.best_params_
Out[344]: {'C': 0.14, 'class_weight': {0: 0.86, 1: 0.83}, 'penalty': 'l2'}

In [345]: gscv_model.best_params_
Out[345]: {'C': 0.15, 'class_weight': {0: 0.85, 1: 0.82}, 'penalty': 'l2'}

In [346]: gscv_model.best_params_
Out[346]: {'C': 0.16, 'class_weight': {0: 0.84, 1: 0.81}, 'penalty': 'l2'}

In [347]: gscv_model.best_params_
Out[347]: {'C': 0.17, 'class_weight': {0: 0.83, 1: 0.8}, 'penalty': 'l2'}

In [348]: gscv_model.best_params_
Out[348]: {'C': 0.18, 'class_weight': {0: 0.82, 1: 0.8}, 'penalty': 'l2'}

```

```
In[243]: {0: 0.19, 1: 0.81}, {0: 0.2, 1: 0.8},
```

```
{0: 0.21, 1: 0.79},
ht_logistic_model = LogisticRegression(penalty='l2', C=
```

```
ht_logistic_model.fit(x_train,y_train)
{0: 0.22, 1: 0.78},
```

```
Out[243]:
```

```
LogisticRegression(C=6.1, class_weight={0:
```

```
0.5, 1: 0.5})
```

```
{0: 0.24, 1: 0.76},
```

```
{0: 0.25, 1: 0.75},
```

```
{0: 0.26, 1: 0.74},
```

```
{0: 0.27, 1: 0.73},
```

```
{0: 0.28, 1: 0.72},
```

```
{0: 0.29, 1: 0.71}, ...],
```

```
'penalty': ['l1',
```

```
'l2']})
```

In [244]:

```
1 # testing data :
2
3 testing_evaluation(ht_logistic_model,x_test,y_test)
```

Confusion Matrix:

```
[[134  16]
 [ 35  46]]
*****
**
```

Accuracy

```
0.7792207792207793
*****
**
```

Classification Report

		precision	recall	f1-score
e	support			
	0	0.79	0.89	0.84
150				
	1	0.74	0.57	0.64
81				
	accuracy			0.78
231				
	macro avg	0.77	0.73	0.74
231				
	weighted avg	0.78	0.78	0.77
231				

In [245]:

```
1 # training data
2
3 training_evaluation(ht_logistic_model,x_train,y_train)
```

Confusion Matrix:

```
[[310  40]
 [ 84 103]]
*****
**
```

Accuracy

```
0.7690875232774674
*****
**
```

Classification Report

		precision	recall	f1-scor
e	support			
	0	0.79	0.89	0.83
350				
	1	0.72	0.55	0.62
187				
	accuracy			0.77
537				
	macro avg	0.75	0.72	0.73
537				
	weighted avg	0.76	0.77	0.76
537				

6.3 Model Training after outlier removing

In [195]:

```
1 logistic_model_OR = LogisticRegression()  
2 logistic_model_OR.fit(x_train, y_train)
```

Out[195]:

LogisticRegression()

In [196]:

```
1 logistic_pred =logistic_model_OR.predict(x_test)  
2 logistic_pred[:5]
```

Out[196]:

array([0, 0, 0, 1, 0], dtype=int64)

6.4 Model Evaluation

In [197]:

```
1 # Testing Data
2 testing_evaluation(logistic_model_OR,x_test,y_test)
```

Confusion Matrix:

```
[[134  16]
 [ 42  39]]
*****
**
```

Accuracy

```
0.7489177489177489
*****
**
```

Classification Report


		precision	recall	f1-scor
e	support			
	0	0.76	0.89	0.82
150				
	1	0.71	0.48	0.57
81				
	accuracy			0.75
231				
	macro avg	0.74	0.69	0.70
231				
	weighted avg	0.74	0.75	0.73
231				

In [198]:

```
1 plot_confusion_matrix(logistic_model_OR,x_test,y_test)
```

Out[198]:

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x198c40af670>
```



In [199]:

```
1 # Training Data
2 training_evaluation(logistic_model_OR,x_train,y_train)
```

Confusion Matrix:

```
[[309  41]
 [ 70 117]]
```

**

Accuracy

0.7932960893854749

**

Classification Report


		precision	recall	f1-scor
e	support			
	0	0.82	0.88	0.85
350				
	1	0.74	0.63	0.68
187				
	accuracy			0.79
537				
	macro avg	0.78	0.75	0.76
537				
	weighted avg	0.79	0.79	0.79
537				

In [200]:

```
1 plot_confusion_matrix(logistic_model_OR,x_train,y_train)
```

Out[200]:

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x198c432d6a0>
```



Pickle File

In [246]:

```
1 with open('Logistic model.pkl', 'wb') as f:  
2     pickle.dump(logistic_model,f)
```

Single user Input Testing

In [247]:

```
1 x_test.iloc[35]
```

Out[247]:

Glucose	141.000000
BloodPressure	72.375171
SkinThickness	29.153420
Insulin	155.548223
BMI	42.400000
DiabetesPedigreeFunction	0.205000
Age	29.000000

Name: 435, dtype: float64

In [248]:

```
1 Glucose = 141.000000
2 BloodPressure = 72.375171
3 SkinThickness = 29.153420
4 Insulin = 155.548223
5 BMI = 42.400000
6 DiabetesPedigreeFunction = 0.205000
7 Age = 29.000000
```

In [249]:

```
1 test_array = np.array([[Glucose, BloodPressure, SkinTh
2 test_array
```

Out[249]:

```
array([[141.      ,  72.375171,  29.15342 ,
        155.548223,  42.4      ,
         0.205    ,  29.      ]])
```

In [250]:

```
1 prediction = logistic_model.predict(test_array)
2 prediction
```

Out[250]:

```
array([1], dtype=int64)
```

In [251]:

```
1 if prediction==1:
2     print("Yes,You are Having Diabetics")
3 else:
4     print("No,You are not Having Diabetics")
```

Yes,You are Having Diabetics

In []:

```
1
```